

CS2102 Lecture 1

Introduction

Database Management System (DBMS)

- What is a DBMS?
 - Software for managing large persistent data
- Advantages of using DBMS
 - Data Independence
 - Efficient Data Access
 - Data Integrity & Security
 - Data Administration
 - Transaction Management
 - Concurrent Access & Crash Recovery
 - Query language

Traditional Data Processing: File Processing Techniques

```
initialize some book-keeping information I
open data file F
while (F is not empty)
    read next record r from F
    if (r satisfies some condition) then
        do something with r
    update I if necessary
do something with I if necessary
close file F
```

Study of DBMS

- Database design
 - How to model the data requirements of applications
 - How to organize data using a DBMS
 - Topics: relational model, ER model, schema refinement
- Database programming
 - How to create, query, and update a database
 - How to specify data constraints
 - How to use SQL in applications
 - Topics: SQL, relational algebra/calculus, stored procedures, triggers
- DBMS implementation
 - How to build a DBMS (Covered in CS3223 & CS4224)

Describing Data in a DBMS

- A DBMS allows users to define and query data in terms of a *data model*
- A **data model** is a collection of concepts for describing data
- A **schema** is a description of the structure of a database using a data model
- A **schema instance** is the content of the database at a particular time

Data Models

- Network Model (e.g., General Electric's IDS (1964))
- Hierarchical Model (e.g., IBM's IMS (1966))
- **Relational Model**
 - **Commercial RDBMS**: IBM DB2, Microsoft SQL Server, Oracle, SAP ASE, etc.
 - **Open-source RDBMS**: MariaDB, MySQL, SQLite, etc.
- Object-oriented Model (e.g., ObjectStore 1988)
- Object-relational Model (e.g., Postgres 1986)
- etc.

Relational Database Systems



(Image: Software Engineering Daily)

Relational Data Model

- Introduced by **Edgar Codd** of IBM Research Laboratory in 1970
- Data is modeled using **relations** (tables with rows & columns)

Students

<i>studentId</i>	<i>name</i>	<i>birthDate</i>	<i>cap</i>
3118	Alice	1999-12-25	3.8
1423	Bob	2000-05-27	4.3
5609	Carol	1999-06-11	4.0

Degree/Arity = Number of columns

Cardinality = Number of rows

- Each relation has a definition called a **relation schema**
 - Schema specifies **attributes** and **data constraints**
 - Data constraints include **domain constraints**

Students (*studentId*: **integer**, *name*: **text**, *birthDate*: **date**, *cap*: **numeric**)

- Each row in a relation is called a **tuple/record**; it has one **component** for each attribute of relation

(1423, 'Bob', 2000-05-27, 4.3)

Relational Data Model (cont.)

- **Domain** - a set of atomic values (e.g., integer, numeric, text)
- Let $\text{domain}(A_i)$ denote the domain of an attribute A_i (set of possible values for A_i)
- Each value of attribute A_i is either a value in $\text{domain}(A_i)$ or **null**
- **null** is a special value used to indicate that the value is either not applicable or unknown
- A relation is a **set of tuples**
 - Consider a relation schema $R(A_1, A_2, \dots, A_n)$ with n attributes A_1, \dots, A_n
 - Each **instance of schema** R is a relation which is a subset of $\{(a_1, a_2, \dots, a_n) \mid a_i \in \text{domain}(A_i) \cup \{\text{null}\}\}$

Relational Data Model (cont.)

- Consider the relation schema **Lectures(course, day, hour)**

- $\text{domain}(\text{course}) = \{\text{'cs101'}, \text{'cs203'}, \text{'cs305'}\}$
- $\text{domain}(\text{day}) = \{1, 2, 3, 4, 5\}$
- $\text{domain}(\text{hour}) = \{8, 10, 12, 14, 16\}$

- Each instance of Lectures is a subset of

$$\{\text{'cs101'}, \text{'cs203'}, \text{'cs305'}, \text{null}\} \times \{1, 2, 3, 4, 5, \text{null}\} \times \{8, 10, 12, 14, 16, \text{null}\}$$

<i>course</i>	<i>day</i>	<i>hour</i>
cs101	1	8
cs101	1	10
cs101	1	12
cs101	1	14
cs101	1	16
cs101	2	8
⋮	⋮	⋮
null	null	16
null	null	null

Relational Data Model (cont.)

- A **relational database schema** consists of a set of relation schemas

Students (*studentId*: **integer**, *name*: **text**, *birthDate*: **date**, *cap*: **numeric**)

Courses (*courseId*: **integer**, *name*: **text**, *credits*: **integer**)

Enrolls (*sid*: **integer**, *cid*: **integer**, *grade*: **numeric**)

- A **relational database** is a collection of tables

Students				Courses		
<i>studentId</i>	<i>name</i>	<i>birthDate</i>	<i>cap</i>	<i>courseId</i>	<i>name</i>	<i>credits</i>
3118	Alice	1999-12-25	3.8	101	Programming in C	5
1423	Bob	2000-05-27	4.3	112	Discrete Mathematics	4
5609	Carol	1999-06-11	4.0	204	Analysis of Algorithms	4
				311	Database Systems	5

Enrolls		
<i>sid</i>	<i>cid</i>	<i>grade</i>
3118	101	5.0
3118	112	4.0
3118	204	3.0
1423	112	4.5
.....

- Relational database schema = relational schemas + data constraints

Relation/Database Schema/Instance

- Relation schema

Students (*studentId*: **integer**, *name*: **text**, *birthDate*: **date**, *cap*: **numeric**)

- Database schema

Students (*studentId*: **integer**, *name*: **text**, *birthDate*: **date**, *cap*: **numeric**)

Courses (*courseId*: **integer**, *name*: **text**, *credits*: **integer**)

Enrolls (*sid*: **integer**, *cid*: **integer**, *grade*: **numeric**)

- Relation (or relation instance)

- Database (or database instance)

Students			
<i>studentId</i>	<i>name</i>	<i>birthDate</i>	<i>cap</i>
3118	Alice	1999-12-25	3.8
1423	Bob	2000-05-27	4.3
5609	Carol	1999-06-11	4.0

Courses		
<i>courseId</i>	<i>name</i>	<i>credits</i>
101	Programming in C	5
112	Discrete Mathematics	4
204	Analysis of Algorithms	4
311	Database Systems	5

Enrolls		
<i>sid</i>	<i>cid</i>	<i>grade</i>
3118	101	5.0
3118	112	4.0
3118	204	3.0
1423	112	4.5
.....

Integrity Constraints (ICs)

- **Integrity constraint**: a condition that restricts the data that can be stored in database instance
 - Specified when schema is defined
 - ICs are checked when relations are updated
- A **legal relation instance** is a relation that satisfies all specified ICs.
- A DBMS enforces ICs - allows only legal instances to be stored

Integrity Constraints (ICs) (cont.)

- Without any additional integrity constraints, each instance of $R(A_1, \dots, A_n) \subseteq \{(a_1, a_2, \dots, a_n) \mid a_i \in \text{domain}(A_i) \cup \{\text{null}\}\}$

Students

<i>studentId</i>	<i>name</i>	<i>birthDate</i>	<i>cap</i>
3118	Alice	1999-12-25	3.8
1423	Bob	2000-05-27	4.3
5609	Carol	1999-06-11	6.5
1423	Dave	2000-10-05	3.7

Courses

<i>courseId</i>	<i>name</i>	<i>credits</i>
101	Programming in C	5
112	Discrete Mathematics	4
204	Analysis of Algorithms	4
null	Compiler Design	4
311	Database Systems	5

Enrolls

<i>sid</i>	<i>cid</i>	<i>grade</i>
3118	101	5.0
3118	112	4.0
3118	202	3.0
1423	112	3.7
5609	101	4.5

Types of Integrity Constraints

- Domain constraints restrict attribute values of relations
- Key constraints
- Foreign key constraints
- Other general constraints

Key Constraints

- A **superkey** is a subset of attributes in a relation that uniquely identifies its tuples
 - No two distinct tuples of a relation have the same values in all attributes of superkey
- **Example:** Which of the following is a superkey for the relation **Students (studentId, name, birthDate, cap)**?
 - {studentId}
 - {name}
 - {birthDate}
 - {cap}
 - {studentId, name}
 - {studentId, birthDate}
 - {studentId, cap}
 - {name, birthDate}
 - {name, cap}
 - {birthDate, cap}
 - {studentId, name, birthDate}
 - {studentId, name, cap}
 - {studentId, birthDate, cap}
 - {name, birthDate, cap}
 - {studentId, name, birthDate, cap}

Key Constraints (cont.)

- A **key** is a superkey that satisfies the additional property:
 - No *proper subset* of the key is a superkey
- Thus, a **key** is a minimal subset of attributes in a relation that uniquely identifies its tuples
- **Example:** Which of the following is a key for the relation Students (studentId, name, birthDate, cap)?
 - {studentId}
 - {name}
 - {birthDate}
 - {cap}
 - {studentId, name}
 - {studentId, birthDate}
 - {studentId, cap}
 - {name, birthDate}
 - {name, cap}
 - {birthDate, cap}
 - {studentId, name, birthDate}
 - {studentId, name, cap}
 - {studentId, birthDate, cap}
 - {name, birthDate, cap}
 - {studentId, name, birthDate, cap}

Key Constraints (cont.)

- Key attribute values cannot be *null*
- A relation could have multiple keys called **candidate keys**
- One of the candidate keys is selected as the **primary key**
- **Example:**
 - Students (studentId, name, email, birthDate)
 - Students has two candidate keys: {*studentId*} and {email}
 - Any one of them could be selected as the primary key

Key Constraints (cont.)

- Consider the relation schema

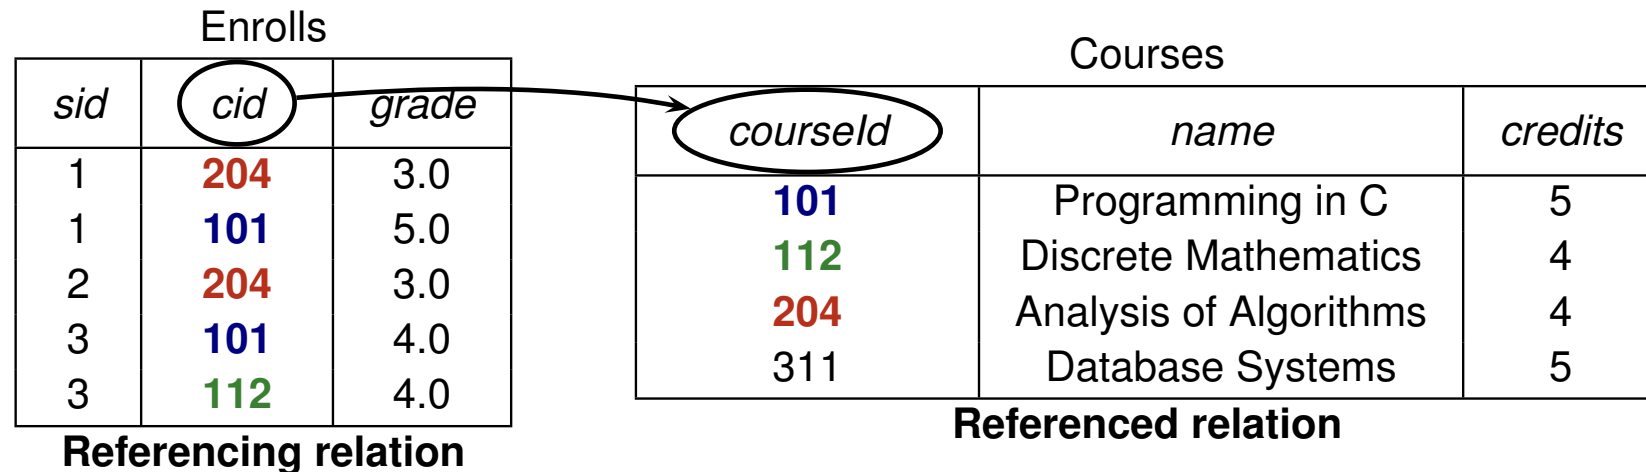
Lectures (cname, pname, day, hour)

- cname** is a course taught by professor **pname** at time given by **day** & **hour**
- Lectures satisfies these constraints:
 - $\text{day} \in \{1, 2, 3, 4, 5\}$
 - $\text{hour} \in \{8, 10, 12, 14, 16\}$
 - At any time, each professor is teaching at most one course
 - Each course is taught by exactly one professor
 - Each course could have multiple lectures

cname	pname	day	hour
cs101	alice	1	10
cs101	alice	3	14
cs200	bob	2	8
ma300	bob	1	10

Foreign Key Constraints

- A subset of attributes in a relation is a **foreign key** if it refers to the primary key of a second relation



- cid* is a foreign key in Enrolls that refers to the primary key *courseid* in Courses
- Foreign key constraint:** each **foreign key** value in **referencing relation** must either (1) appear as **primary key** value in **referenced relation** or (2) be a null value

Foreign Key Constraints (cont.)

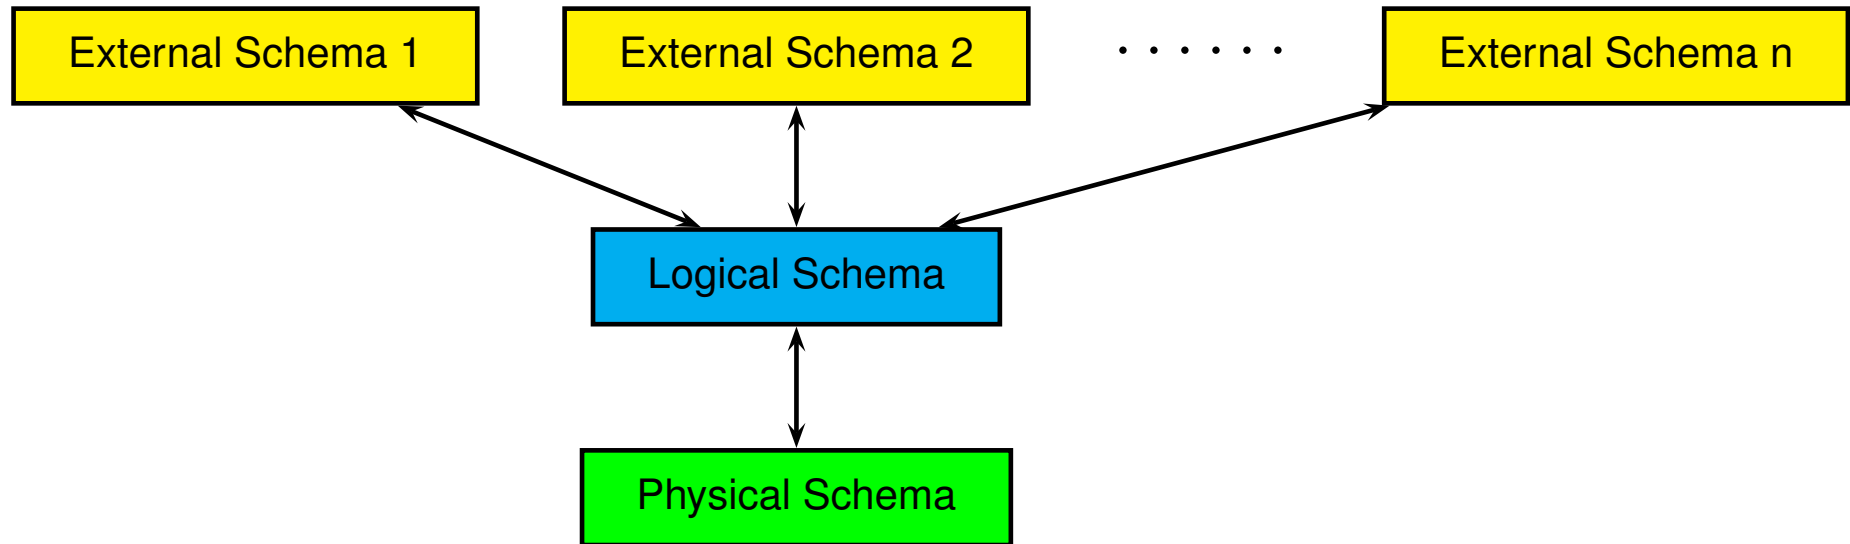
- The referencing & referenced relations could be the same relation
- **Example:** Each employee has at most one manager

Employees

<i>eid</i>	<i>ename</i>	<i>managerid</i>
001	Alice	<i>null</i>
002	Bob	001
003	Carol	001
007	Dave	<i>null</i>
008	Eve	007

- Constraints on Employees table:
 - *eid* is the primary key
 - *managerid* is a foreign key that refers to *eid*
- Foreign key constraints are also known as **referential integrity constraints**

Levels of Data Abstraction



- Data in DBMS is described at three levels of abstractions
- **Logical Schema** - logical structure of data in DBMS
- **Physical Schema** - how the data described by logical schema is physically organized in DBMS
- **External Schema** - A customized view of logical schema for a group of users or an individual user

External Schema Example

- Consider the following logical database schema:

Students (*studentId, sname, birthDate, cap*)

Profs (*profId, pname, email, office*)

Courses (*courseId, cname, credits, profId, lectureTime*)

Enrolls (*sid, cid, grade*)

- External schema for Alice:

CourseEnrollment (*cname, pname, lectureTime, totalEnrollment*)

- External schema for Bob:

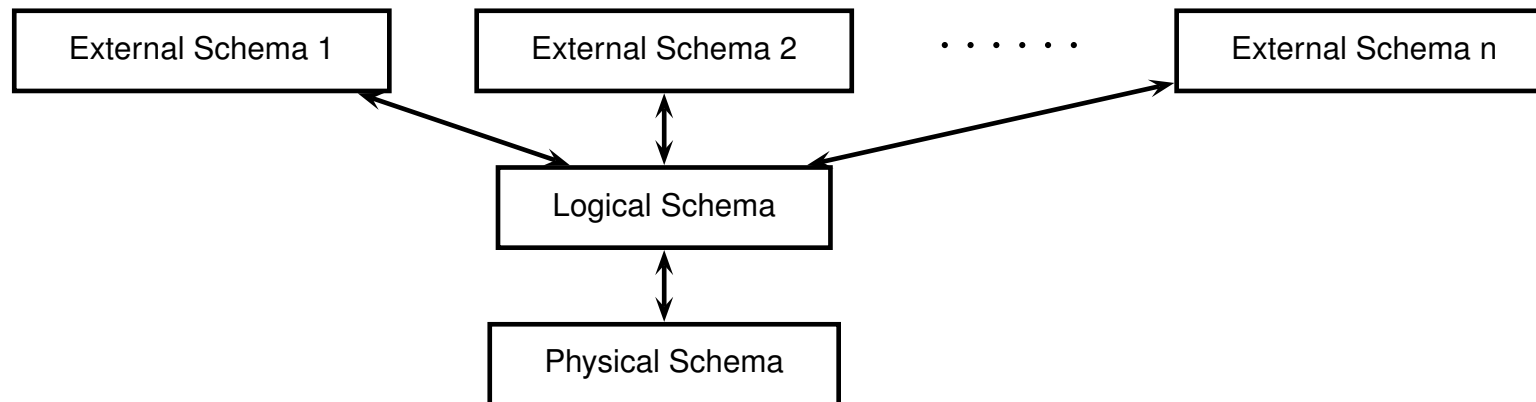
StudentInfo (*studentId, sname*)

CourseInfo (*courseId, cname, credits, profId, lectureTime*)

EnrollInfo (*sid, cid, cname, pname, lectureTime*)

Data Independence

- Insulate users/applications from changes in how data is structured and stored
- Data independence is achieved via the three levels of abstraction



- **Physical data independence** - protection from changes in physical schema
- **Logical data independence** - protection from changes in logical schema
- Data independence is an important advantage of using DBMS!

Transactions

- Abstraction for representing a logical unit of work
- **ACID Properties**
 - **Atomicity**: Either all the effects of a transaction are reflected in the database or none are
 - **Consistency**: The execution of a transaction in isolation preserves the consistency of the database
 - **Isolation**: The execution of a transaction is isolated from the effects of other concurrent transaction executions
 - **Durability**: The effects of a committed transaction persists in the database even in the presence of system failures

Transaction Example

Transfer(X, Y, amount)

```
fromBal := read balance from X's account
if fromBal  $\geq$  amount then
    toBal := read balance from Y's account
    update Y's balance to toBal + amount
    update X's balance to fromBal - amount
end if
```

Serial Transaction Executions

Two possible serial executions of `Transfer(1,2,100)` & `Transfer(2,1,100)`

(1): `fromBal := read 1's balance`
`toBal := read 2's balance`
`Update 2's balance to toBal + 100`
`Update 1's balance to fromBal - 100`
`fromBal := read 2's balance`
`toBal := read 1's balance`
`Update 1's balance to toBal + 100`
`Update 2's balance to fromBal - 100`

(2): `fromBal := read 2's balance`
`toBal := read 1's balance`
`Update 1's balance to toBal + 100`
`Update 2's balance to fromBal - 100`
`fromBal := read 1's balance`
`toBal := read 2's balance`
`Update 2's balance to toBal + 100`
`Update 1's balance to fromBal - 100`

Concurrent Transaction Executions

A concurrent execution of $\text{Transfer}(1,2,100)$ & $\text{Transfer}(2,1,100)$

fromBal := read 1's balance

toBal := read 2's balance

Update 2's balance to toBal + 100

fromBal := read 2's balance

toBal := read 1's balance

Update 1's balance to fromBal - 100

Update 1's balance to toBal + 100

Update 2's balance to fromBal - 100

Summary

- DBMS used to store, update, and query data
- **Relational data model**
 - Tabular representation of data
 - Integrity constraints specify restrictions on data based on application semantics
- Levels of data abstraction provide data independence
- Transactions simplify application development